

Random Forests and Boosting

Christoph Kern

Mannheim Machine Learning Modules

c.kern@uni-mannheim.de

February 6 and 7, 2018



Outline

- 1 Introduction
- 2 Bagging
- 3 Random Forests
 - Growing a Forest
 - Interpretation
- 4 Boosting
 - AdaBoost
 - GBM
- 5 Summary
- 6 Software Resources
- 7 References

Introduction

Some limitations of (single) trees

- Difficulties in modeling additive structures
- Lack of smoothness of prediction surface
- High variance / **instability** due to hierarchical splitting process

→ **Ensemble methods**

- Address instability via combining multiple prediction models
- Combine diverse models into a more robust ensemble

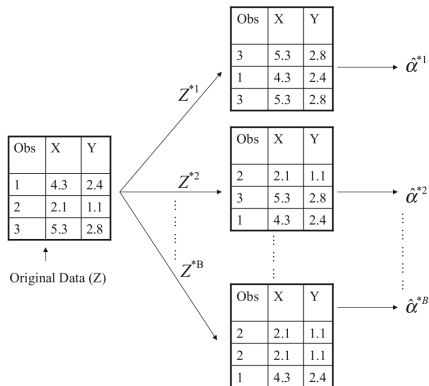
How to construct ensembles?

- Combine models based on different methods
 - Stacking: Build a meta-model that uses (multiple) predictions as input
- Apply one method with different tuning parameter settings
- Combine models with different features
- Use one method with different subsets of the data
 - **Bagging**: Can be applied to different **base learners** (e.g. CART)

- 1 Introduction
- 2 Bagging**
- 3 Random Forests
 - Growing a Forest
 - Interpretation
- 4 Boosting
 - AdaBoost
 - GBM
- 5 Summary
- 6 Software Resources
- 7 References

Bagging

Figure: Bootstrap process



James et al. (2013)

Bootstrap: Sampling B samples of size n with replacement from original data set




Applications

- Estimate the variability of model parameters
 - e.g. standard errors of regression coefficients
- Estimate test error with training data
 - Fit model on bootstrap samples and predict original training set
- Construct an ensemble of learners for prediction
 - **Bagging**: Bootstrap Aggregating
 - Train prediction models on bootstrap samples

Algorithm 1: Bagging Trees

```
1 Set number of trees  $B$ ;  
2 Define stopping criteria;  
3 for  $b = 1$  to  $B$  do  
4   | draw a bootstrap sample from the training data;  
5   | assign sampled data to root node;  
6   | if stopping criterion is reached then  
7   |   | end splitting;  
8   | else  
9   |   | find the optimal split point among the predictor space;  
10  |   | split node into two subnodes at this split point;  
11  |   | for each node of the current tree do  
12  |   |   | continue tree growing process;  
13  |   | end  
14  | end  
15 end
```

Figure: Bagging Trees

(a) $b = 1$ (b) $b = 2$ (c) $b = 3$ 

- 1 Introduction
- 2 Bagging
- 3 Random Forests**
 - Growing a Forest
 - Interpretation
- 4 Boosting
 - AdaBoost
 - GBM
- 5 Summary
- 6 Software Resources
- 7 References

Random Forests

From Bagging to Random Forests

Variance of an average of B i.i.d. random variables

$$\frac{1}{B}\sigma^2$$

→ Bagging: Averaging over B trees decreases variance

Variance of an average of B i.d. random variables with $\rho > 0$

$$\rho\sigma^2 + \frac{1-\rho}{B}\sigma^2$$

→ **Random Forests:** Averaging over B trees with m out of p predictors per split decreases variance and decorrelates trees

Random Forests

From Bagging to Random Forests

Variance of an average of B i.i.d. random variables

$$\frac{1}{B}\sigma^2$$

→ Bagging: Averaging over B trees decreases variance

Variance of an average of B i.d. random variables with $\rho > 0$

$$\rho\sigma^2 + \frac{1-\rho}{B}\sigma^2$$

→ **Random Forests:** Averaging over B trees with m out of p predictors per split decreases variance and decorrelates trees

The Random Forest trick

- Randomization with respect to rows *and* columns
- Weaker predictors have more of a chance
- Results in diverse and *decorrelated* trees

Can be taken one step further...

- 1 Draw a random sample m from the p predictors (w/o Bootstrapping)
- 2 Draw random split(s) per feature
- 3 Split node using the best of these random splits

→ Extremely Randomized Trees (Geurts et al. 2006)

Growing a Forest

Algorithm 2: Grow a Random Forest

```
1 Set number of trees  $B$ ;  
2 Set predictor subset size  $m$ ;  
3 Define stopping criteria;  
4 for  $b = 1$  to  $B$  do  
5   draw a bootstrap sample from the training data;  
6   assign sampled data to root node;  
7   if stopping criterion is reached then  
8     end splitting;  
9   else  
10    draw a random sample  $m$  from the  $p$  predictors;  
11    find the optimal split point among  $m$ ;  
12    split node into two subnodes at this split point;  
13    for each node of the current tree do  
14      continue tree growing process;  
15    end  
16  end  
17 end
```

A Random Forest

$$\{T_b\}_1^B$$

consists of a set of $b = 1, 2, \dots, B$ trees which can be used for prediction by...

- Regression

- ...averaging predictions over all trees

- $\hat{f}_{rf}^B(x) = \frac{1}{B} \sum_{b=1}^B T_b(x)$

- Classification

- ...using most commonly occurring class among all trees

- $\hat{C}_{rf}^B(x) = \text{majority vote}\{\hat{C}_b(x)\}_1^B$

Observations in each bootstrap sample

$$\begin{aligned} P(\text{obs } i \in \text{sample } b) &= 1 - \left(1 - \frac{1}{n}\right)^n \\ &\approx 1 - e^{-1} \\ &= 0.632 \end{aligned}$$

Out-of-bag (OOB) error

- Sampling with replacement leads to models based on subsets of the data
- Unused (OOB) observations can be used for test error estimation
 - 1 Generate predictions for case i using models where i was OOB
 - 2 Average predictions for i and estimate test error
 - 3 Compute OOB error over all cases

Tuning Random Forests

- Predictor subset size m out of p
 - Most important tuning parameter in RF
 - Starting value; $m = \sqrt{p}$ (classification), $m = p/3$ (regression)
 - Can be chosen using OOB errors based on different m
- Optional: Number of trees
 - sufficiently high (e.g. 500)
- Optional: Node size (number of observations in terminal nodes)
 - sufficiently low (e.g. 5)

Interpretation

Interpreting Random Forests

- Inspect each tree of the forest
 - Inefficient for 500+ trees
- Variable importance
 - Summary of “effect size”
- Partial dependence plots
 - Graphical representation of “effect structure”
 - Outlook: ICE plots (Goldstein et al. 2014)

Variable importance with CART

$$\mathcal{I}_\ell^2(T) = \sum_{t=1}^{J-1} \hat{i}_t^2 I(v(t) = \ell)$$

- Sum of squared improvements \hat{i}^2 over all internal nodes with predictor X_ℓ
 - Regression: Overall reduction in RSS caused by X_ℓ
 - Classification: Overall reduction of impurity caused by X_ℓ

Importance with Random Forests

$$\mathcal{I}_\ell^2 = \frac{1}{M} \sum_{m=1}^M \mathcal{I}_\ell^2(T_m)$$

- Average improvement caused by predictor X_ℓ over all trees


Partial dependence plots

$$\tilde{f}(x) = \frac{1}{n} \sum_{i=1}^n f(x, x_{iC})$$


- Goal: Plot results from “black box” learning methods (e.g. RF)
- Compute $\tilde{f}(x)$ over the range of x while averaging the effects of the remaining predictors x_{iC}
- Generate artificial datasets by fixing x for all cases
 - Regression: Averaging over $f(x, x_{iC})$ for each value of x
 - Classification: Averaging over $\text{logit}(p)$ for each value of x

Figure: Partial dependence plots

(a) CART



(b) Random Forest



- 1 Introduction
- 2 Bagging
- 3 Random Forests
 - Growing a Forest
 - Interpretation
- 4 Boosting**
 - AdaBoost
 - GBM
- 5 Summary
- 6 Software Resources
- 7 References

Boosting

Boosting

- Class of ensemble methods which combine **sequential** prediction models
- Adaptive approach with focus on “difficult observations”
- Different flavors exist
 - AdaBoost
 - Gradient Boosting Machines (GBM)
 - ...
- Can be applied to different (weak) base learners
 - Boosting trees
 - ...

AdaBoost

AdaBoost


- Algorithm for classification problems ($Y \in \{-1, 1\}$)
- Estimate a sequence of classifiers using reweighted data
- AdaBoost process
 - 1 Fit classifier $G_m(x)$ to weighted data (initial weights $w_i = \frac{1}{n}$)
 - 2 Compute the misclassification rate

$$\text{err}_m = \frac{\sum_{i=1}^n w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^n w_i}$$


- 3 Compute the classifier weight $\alpha_m = \log((1 - \text{err}_m)/\text{err}_m)$
 - 4 Recalculate weights $w_i = w_i \exp(\alpha_m I(y_i \neq G_m(x_i)))$
- Majority vote classification: $G(x) = \text{sign} \left[\sum_{m=1}^M \alpha_m G_m(x) \right]$

Figure: (Ada)Boosting stumps (example)¹


(a) Step 1: $\alpha_1 = 0.42$



(b) Step 2: $\alpha_2 = 0.65$




(c) Step 3: $\alpha_3 = 0.92$



¹http://www.ccs.neu.edu/home/vip/teach/MLcourse/4_boosting/slides/boosting.pdf

Figure: Step 4: Combine models



GBM

Gradient Boosting Machines (GBM)

- General approach to sequential learning
- Applicable with various loss functions
- Boosting trees
 - 1 Initialize model (with a constant $f_0(x)$)
 - 2 Compute pseudo-residuals based on current model

$$r_{im} = - \left[\frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \right]_{f=f_{m-1}}$$

- 3 Fit a regression tree to the pseudo-residuals
 - 4 Compute $\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(y_i, f_{m-1}(x_i) + \gamma)$
 - 5 Update the current model: $f_m(x) = f_{m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$
- Output $\hat{f}(x) = f_M(x)$

→ Analogue to steepest descent

Table: GBM components for different loss functions

Setting	Loss function	r_i	$f_0(x)$
Regression	$\frac{1}{2}(y_i - f(x_i))^2$	$y_i - f(x_i)$	mean(y_i)
Regression	$ y_i - f(x_i) $	$\text{sign}(y_i - f(x_i))$	median(y_i)
Classification	Deviance	$I(y_i = G_k) - p_k(x_i)$	prior p's

Shrinkage

- Additional tweak in Gradient boosting
- Slow down learning rate to avoid overfitting
- Learning rate is controlled by λ
 - $f_m(x) = f_{m-1}(x) + \lambda \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$

Subsampling

- Optional add-on in Gradient boosting
- Use a random sample (w/o replacement) of pseudo-residuals in each step
- Can be introduced to improve performance and speed
 - “Stochastic gradient boosting”

Algorithm 3: Gradient Boosting for regression

```
1 Set number of trees  $M$ ;  
2 Set interaction depth  $D$ ;  
3 Set shrinkage parameter  $\lambda$ ;  
4 Use  $\bar{y}$  as initial prediction;  
5 for  $b=1$  to  $B$  do  
6   | compute residuals based on current predictions;  
7   | assign data to root node, using the residuals as the outcome;  
8   | while current tree depth  $< D$  do  
9   |   | tree growing process;  
10  | end  
11  | compute the predicted values of the current tree;  
12  | add the shrinked new predictions to the previous predicted values;  
13 end
```

Table: Gradient Boosting with 5 obs and 2 x's (example)²

ID	x_1	x_2	y	$f_0(x)$
1	0	0	1	1.2
2	0	2	3	1.2
3	1	2	2	1.2
4	2	3	0	1.2
5	0	1	0	1.2

²Schonlau, M. (2015). Statistical Learning with Boosting. ESRA short course.

Table: Step 1: Split $x_2 > 2.5$

ID	x_1	x_2	y	$f_0(x)$	r_{i1}	γ_{j1}	$f_1(x)$
1	0	0	1	1.2	-0.2	0.3	1.5
2	0	2	3	1.2	1.8	0.3	1.5
3	1	2	2	1.2	0.8	0.3	1.5
4	2	3	0	1.2	-1.2	-1.2	0
5	0	1	0	1.2	-1.2	0.3	1.5

Table: Step 2: Split $x_2 < 1.5$

ID	x_1	x_2	y	$f_0(x)$	$f_1(x)$	r_{i2}	γ_{j2}	$f_2(x)$
1	0	0	1	1.2	1.5	-0.5	-1	0.5
2	0	2	3	1.2	1.5	1.5	0.66	2.166
3	1	2	2	1.2	1.5	0.5	0.66	2.166
4	2	3	0	1.2	0	0	0.66	0.66
5	0	1	0	1.2	1.5	-1.5	-1	0.5

Tuning Gradient Boosting Machines

- Number of trees M
 - Number of “iterations”
 - Overfitting can occur for large M
- Interaction depth D
 - Number of splits for each tree
 - Boosting stumps: $D = 1$
- Shrinkage parameter λ
 - e.g. $\lambda = 0.01$, $\lambda = 0.001$
 - Smaller λ needs larger M
- ...

Summary

- Ensemble methods combine multiple models to stabilize predictions
- RF and Boosting are competitive “general purpose” approaches
- A lot of different flavors exist
- Algorithms typically compared in a large train and tune loop
- Drawbacks: Lower interpretability and higher computational costs

Software Resources

Resources for R

- Standard package to grow RFs: `randomForest`
- Extremely Randomized Trees: `extraTrees`
- Gradient Boosting: `gbm`
- Extreme Gradient Boosting: `xgboost`
- Visualization
 - Partial Dependence Plots: `pdp`
 - Plot model surfaces (also PDPs): `plotmo`

References

- Berk, R. A. (2006). An Introduction to Ensemble Methods for Data Analysis. *Sociological Methods & Research*, 34(3), 263–295.
- Biau, G., Scornet, E. (2015). *A Random Forest Guided Tour*. arXiv: 1511.05741.
- Friedman, J. (2001). Greedy Function Approximation: A Gradient Boosting Machine. *The Annals of Statistics*, 29(5), 1189–1232.
- Geurts, P., Ernst, D., Wehenkel, L. (2006). Extremely Randomized Trees. *Machine Learning* 63(1), 3–42.
- Goldstein, A., Kapelner, A., Bleich, J., Pitkin, E. (2014). *Peeking Inside the Black Box: Visualizing Statistical Learning with Plots of Individual Conditional Expectation*. arXiv: 1309.6392v2.
- James, G., Witten, D., Hastie, T., Tibshirani, R. (2013). *An Introduction to Statistical Learning*. New York, NY: Springer.